Anshu Kumar Sharma · 22322008

Read the script, *don't watch the show.*

Today an editor sits through **all 50 episodes** to pick promo moments, and Kuku adds a show every week. Every episode already arrives as a **timestamped screenplay**, so I let software read it instead and hand the editor a **ranked list of 20 clips** with exact in/out times. The editor's job shrinks to approving.

A THE BRIEF, TURNED INTO CHECKS

20–60 s	Windows are only ever built inside the range; a model edit that leaves it is refused.	CODE
Hook ≤ 3 s	A candidate needs dramatic evidence in its opening 3 s; the judge then scores it 1–5.	CODE
Unresolved cut	Ends on a beat boundary; the judge rates how open the last note is.	AI
Works cold	The judge sees only the clip plus a 3-beat lead-in, like a stranger would.	AI

B ONE RUN, ALL 50 EPISODES

Each step throws most of the work away

50	episodes in	.txt screenplays
3,727	timed beats	parse · 0 warnings
200	candidates	rules · 4/ep
152	kept by judge	Gemini
20	on the slate	rank

Rules filter cheaply, so the bill is fixed before a run: **200 calls per show**. Ranking is plain code with readable weights that an editor can argue with.

C THE ONE DESIGN RULE

The model never writes a timestamp

Parser CODE

Turns every line, action, camera and sound cue into a beat with an id and exact times.

```
E03.B064  77.2–78.9s
E03.B065  79.0–81.0s
…
```

→

Judge AI

Reads the beats as text. May only answer with ids from the clip it was shown.

```
{"start_beat":
"E03.B064",
 "end_beat":
"E03.B112",
 "hook_strength": 5, …}
```

→

Validator CODE

Id in the clip? In order? Still 20–60 s? If not, keep the original window and log a flag.

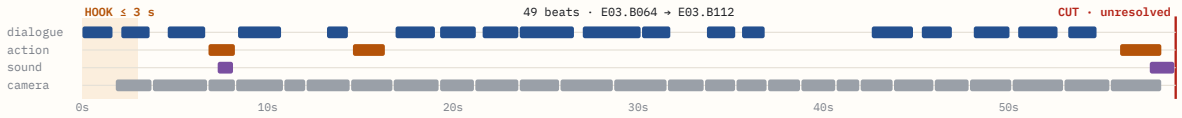
```
ids → times
01:17.20 → 02:16.20
```

0
ways to output a made-up time: times only come from the parser

118/200
clips the judge tightened, so the id mechanism is really used

8
edits that broke 20 s, caught and reverted, all logged in `run_report.json`

D WHAT THE EDITOR RECEIVES · PROMO #1 · EP 03 · 01:17.20–02:16.20 · 59 S



HOOK A scarred stranger walks into a chapel and drops a royal guard with one strike.

CUT ON He kneels to propose to a woman he has just met; the congregation freezes.

SCORES `hook 5 · cliffhanger 5 · standalone 5 · spoiler 1 → 4.40`

E WHY I PICKED THIS PROBLEM

Real input. It was the only option with an actual data pack, so results can be checked, not imagined.

Recurring pain. A new show every week means the saving compounds; it is not a one-off.

A testable definition. The brief says what a promo is (length, hook, open ending), so I could score against it.

Output: `promos.md` · `promos.csv` · `promos.json`, plus every rejected clip with its reason.

Run: all 50 episodes, gemini-3.1-flash-lite, 980 s. Every figure is read from outputs/ and run_report.json.

Live demo runs in the browser · bring any model + key

Is it any good? *I checked 30 clips by hand.*

I judged each clip from the beats in the screenplay, never from the model's own summary of it. Then I fixed what the labels exposed and measured again.

F HAND LABELS AGAINST THE BRIEF

SLICE	N	CLIPS	READING
Top 10	10	6 post it · 4 borderline · 0 bad	**60%** clean · 90% usable
Whole slate	18	8 post it · 7 borderline · 3 bad	**44%** clean · 89% usable
Rejected	12	3 borderline · 9 bad	**67%** rightly dropped

post it · borderline · bad

eval/eval_report.md

G FIX & RE-MEASURE

The labels exposed three bugs in the rules: start-of-scene position counted as a hook, and `crack` / `fire` matched a cosy "crackling hearth fire".

MEASURE	BEFORE	AFTER
Top-10 clean	~~30%~~	60%
Slate usable	~~65%~~	89%
Rejects right	~~77%~~	67%

The cost: a stricter gate means the judge's rejects are closer calls. I report it rather than hide it.

H MECHANICAL CHECKS

100% timecodes land on real beat boundaries

100% promos inside 20–60 s

200/200 candidates judged, 0 API failures

0 parse warnings across 3,727 beats

`eval/verify_outputs.py` re-parses the scripts on its own and checks every output against them.

I WHERE IT STILL PICKS BADLY

6/30 **Soft endings.** The judge gives a 4/5 cliffhanger to clips that end on a resolution, e.g. a flirt. Absolute 1–5 scores are the weak spot.

4/12 **Too strict at the gate.** Some rejects were postable, one ending on the single word "Arsenic."

198 **Spoiler score barely varies.** 198 of 200 clips got 1 or 2, so it hardly moves the ranking.

— **Reads words, not pictures.** A striking silent shot is underrated. Any text-only method shares this limit.

J SPEED, COST, SCALE

Same scores, 100× the wait. A reasoning model and Flash Lite rated the same clip identically.



Calls per show	200
Tokens per show	~230k
Free tier (15/min, 500/day)	~16 min
Paid tier	< 1 min

Cost grows with episode count, not episode length, so the budget is known before a run.

K TRY IT

Web: drop in screenplays, pick any model (Gemini, OpenAI, Claude, Groq…) and paste a key. It runs in the browser; the key never touches a server.

```
# or locally, any provider
python run.py --input <pack> --out outputs/
  --provider gemini|openai|anthropic|groq
# free re-rank, no API calls
  --reuse-judgements outputs/judgements.json
```

L GAPS IN THE BRIEF I HAD TO FILL

1. Script times match the cut video 1:1.
2. 20 and 60 s are both allowed.
3. Clips start and end on whole beats.
4. "Makes sense alone" is judged on the clip only.
5. Joining two segments is allowed; one stretch preferred.
6. Spoiler risk grows with episode number.

M BUGS MY AI ASSISTANTS INTRODUCED

- **Retired model id → 404.** Now I list models from the API, not from memory.
- **"Beats are always in time order."** False in 8 places; a parser test caught it.
- **Hearth fire read as violence.** Hand labels caught it; fixed in G.
- **A run hung for 2 hours.** The SDK has no default timeout; I added one.

N WHAT I'D BUILD NEXT

Pairwise judging instead of lone 1–5 scores; it targets both failure modes.

Learned weights from real watch-through data.

Keyframes for top picks, to catch silent moments.